

BIG DATA [MODULE 2]

Cloud computing

whoami

Matteo Francia

- Email: m.francia@unibo.it
- Research fellow @ UniBO

Research topics

- Big data / database
- Geo-spatial analytics

Thesis proposals

- <https://big.csr.unibo.it/teaching/>

<https://big.csr.unibo.it/>



Exam

We will see how (big) data pipelines can be done in the cloud

- We will cover some "theory" as well as get our hands dirty
- This module is part of the oral examination
- (If agreed) The project can be extended to the cloud to get more points
 - E.g., rethink a part of the project on cloud

Roadmap

Introduction to cloud and cloud service providers

Cloud services (in AWS)

Hands-on cloud services

- Orchestrating (small) data pipelines
- Migrating a cluster to the cloud
- Migration pricing

Set up

You (will) have an AWS Educate account

- A coupon of 150\$ to test AWS cloud services

The content of these slides refers to this repo

- <https://github.com/w4bo/bigdata-aws/>
- Options
 - Work on your pc: check the `README.md` to install the necessary software tools
 - Mainly: git, IntelliJ IDEA (Community Edition), docker, AWS CLI, AWS SAM CLI
 - Download a pre-configured (VMware) VM with the installed software

<https://aws.amazon.com/education/awseducate/>

Set up

Download a pre-configured (VMware) VM with the installed software

1. Connect to a PC-lab using Guacamole (click here <https://csi-rlab.campusfc.unibo.it/>)
 - a) Enter your credentials
 - b) Choose `LabCEZ`
 - c) Click on the proper lab
2. Once logged in, download the VM from <http://big.csr.unibo.it/downloads/ubuntu-64-bit.zip>
3. Unzip the content and launch `VMware Workstation Pro` (from the desktop)
4. `File` > `Open...` > `path/to/the/unzipped/vm/*.vmx` file
5. Launch the VM (user: `bigdata`, pwd: `bigdata`)

So far

You have practiced with **on-premises** solutions

- You were given a working hardware cluster
- ... to deploy software applications on Hadoop-based stack

Let us guess, how would you start from scratch?

- How would you do that?
- How much time would it take?

So far

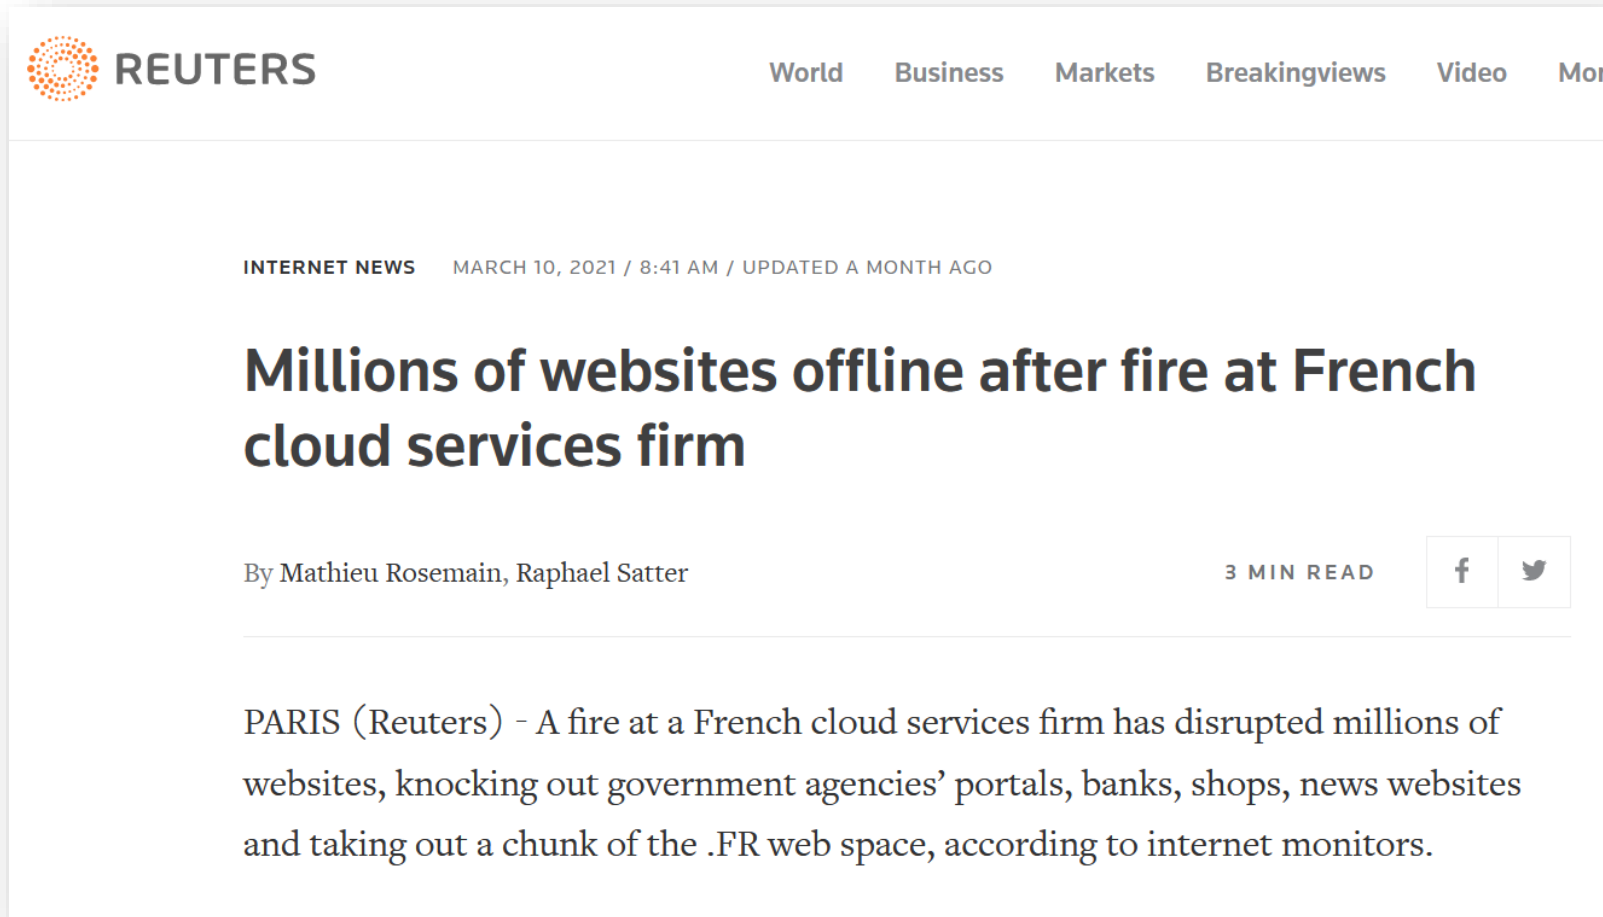
No easy answers

Big-data (distributed) architectures require a lot of skills

- **Installation:** how do I set up a new machine?
- **Networking:** how do I cable dozens of machines?
- **Management:** how do I replace a broken disk?
- **Upgrade:** how do I extend the cluster with new services/machines?
- (energy and cooling, software licenses, insurance...)

<https://aws.amazon.com/compliance/data-center/data-centers/>

So far



<https://www.reuters.com/article/us-france-ovh-fire-idUSKBN2B20NU>

So far

Technological perspective

- How do we configure a distributed environment?
 - How do we set up independent services?
 - How do we integrate such services?
 - How do control resource accesses (e.g., storage)?
- How do we orchestrate data flows?
- It depends on your (team) skills (not only software engineering)

Business perspective

- No free lunch, each choice has cost/benefit
- How much time does it take to master a technology?
- How many people do I need?

So far

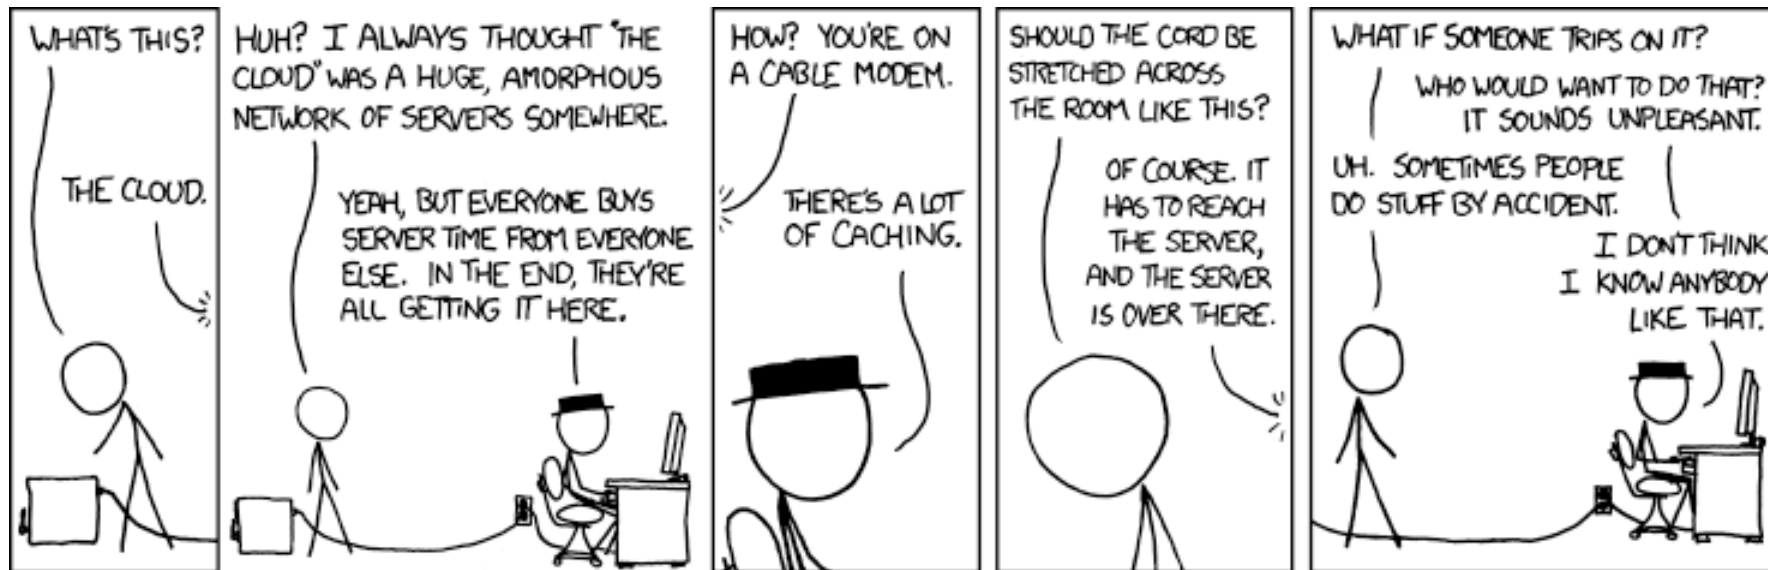
Can we afford to spend resources on tasks that are not mission oriented?

- Mission: a statement used by a company to explain its purpose(s)

How can I build a working application/data platform?

BIG DATA




Why going cloud?



<https://xkcd.com/908/>



Teaching material

Books



Generic —————> Specific

Web content



Why going cloud?

Cloud computing (National Institute of Standards and Technology)

*“A model for enabling **ubiquitous, convenient, on-demand** network access to a **shared pool** of configurable computing resources (e.g., networks, servers, storage, services) that can be rapidly provisioned and released with **minimal management effort** or service provider interaction.”*

- On-demand self-service (consume services when you want)
- Broad network access (consume services from anywhere)
- Resource pooling (infrastructure, virtual platforms, and applications)
- Rapid elasticity (enable horizontal scalability)
- Measured service (pay for the service you consume as you consume)

Why going cloud?

Goal: adjusts capacity to have predictable performance at the lowest cost

Scalability that is not possible on premises

- Scale from one to thousands of servers

Elasticity

- Automatically scale resources in response to run-time conditions
- Core justification for the adoption of cloud

Why going cloud

Hardware scalability

- No longer think about rack space, switches, and power supplies, etc.


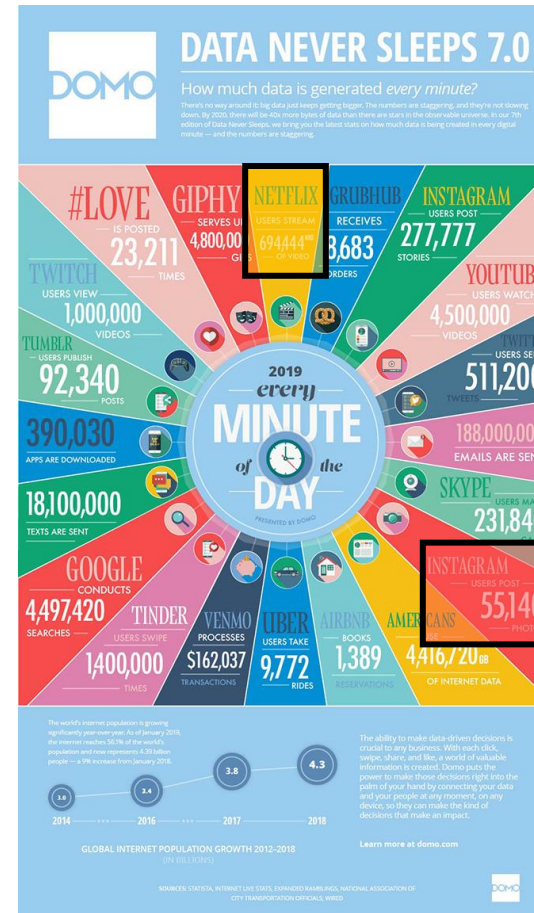

Grow storage from GBs to PBs

- One hundred of our 10TB Enterprise Capacity 3.5 HDD hard drives



<https://blog.seagate.com/business/linus-tech-tips-want-petabyte-system/>

Why going cloud?



<https://www.domo.com/learn/data-never-sleeps-8>

Why going cloud?

Resource pooling

- Enable a resource to serve different consumers
- Dynamically reassigned according to demands
- Economy of scale

Reliability

- Built to handle failures
- Fault-tolerant or highly available

Worldwide deployment

- Deploy applications as close to customers as possible
 - E.g., to reduce network latency
- Improve data locality
- Compliant to privacy regulations (e.g., GDPR)

Why going cloud?

User perspective

- Eliminate repetitive tasks to focus on strategic ones
- Abstract the underlying architecture
- Adapt infrastructure to requirements, create (test) environments on demand

Service **integration**

- Do not reinvent the wheel
- Use services that solve common problems (e.g., load balancing, queuing)

Service integration and abstraction are drivers of change

- From databases to data platforms
- From on-premises hardware to serverless architectures
- From custom to standard data pipelines

BIG DATA

From databases to data platforms

Data platform

Companies are collecting tons of data to enable advanced analytics

- Data are more and more heterogeneous and complex
- Raw data are difficult to obtain, interpret, describe, and maintain
- There is a need for describing/curating the data to make them consumable

Databases/warehouses are no longer ideal data hubs for integration/analysis

- Getting value from data is not only a matter of storage
- Need integrated and multilevel analytical skills and techniques

Data platform

Database

"A database is a structured and persistent collection of information about some aspect of the real world organized and stored in a way that facilitates efficient retrieval and modification. The structure of a database is determined by an abstract data model. Primarily, it is this structure that differentiates a database from a data file. The most popular data model is relational that represents data as a set of tables."

Data Warehouse

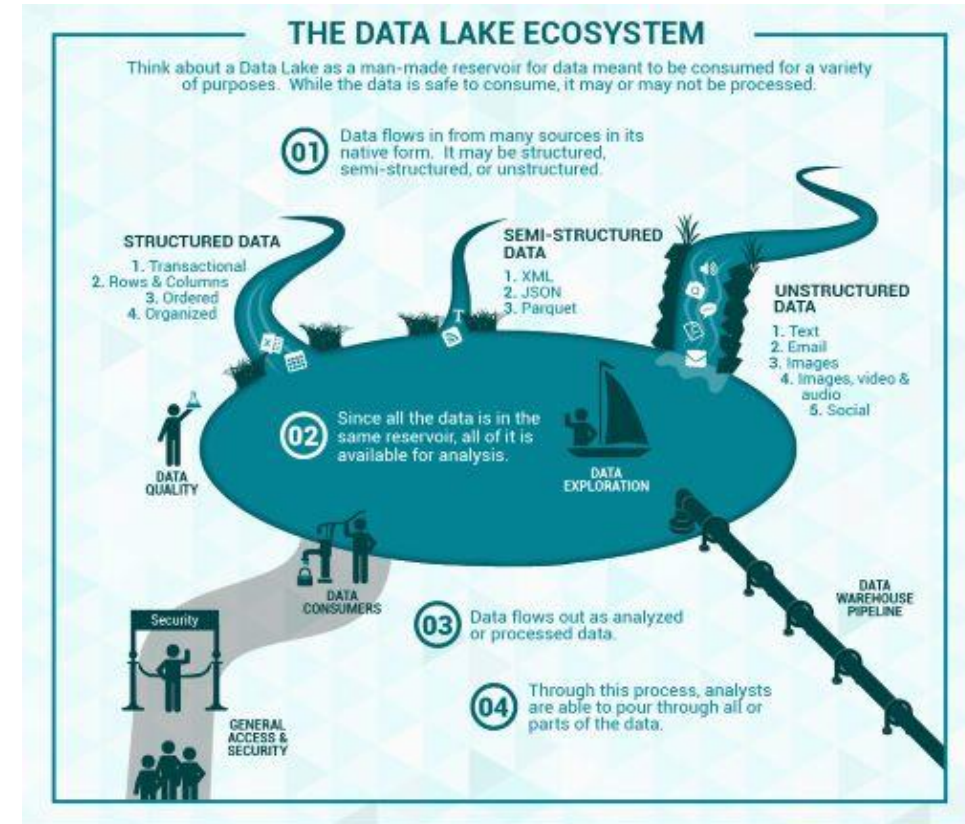
"A collection of data that supports decision-making processes. It provides the following features: subject-oriented, integrated and consistent, not volatile."

Özsu M.T. (2018) Database. In: Encyclopedia of Database Systems. Springer, New York, NY. https://doi.org/10.1007/978-1-4614-8265-9_80734
Matteo Golfarelli and Stefano Rizzi. *Data warehouse design: Modern principles and methodologies*. McGraw-Hill, Inc., 2009.

Data platform

Data lake

Couto et al.: “A DL is a **central repository** system for **storage, processing, and analysis of raw data**, in which the data is kept in its **original format and is processed to be queried only when needed**. It can **store a varied amount of formats** in big data ecosystems, from unstructured, semi-structured, to structured data sources”



Couto, Julia, et al. "A Mapping Study about Data Lakes: An Improved Definition and Possible Architectures." *SEKE*. 2019.
<https://dunnsolutions.com/business-analytics/big-data-analytics/data-lake-consulting>

Data platform

Data lakes have increasingly taken the role of data hubs

- Eliminate up-front costs of ingestion since data are stored in original format
- Once in DL, data are available for analysis by everyone in the organization

Drawing a sharp line between storage/computation/analysis is hard

- Is a database just storage?
- What about SQL/OLAP?

Blurring of the architectural borderlines

- DL is often replaced by “data platform” or “data ecosystem”
- Encompass systems supporting data-intensive storage, computation, analysis

Data platform

Data platform (e.g., on Google Cloud and Amazon AWS)

- Rationale: relieve users from complexity of administration and provision
 - Not only technological skills, but also privacy, access control, etc.
 - Only focus on functional aspects

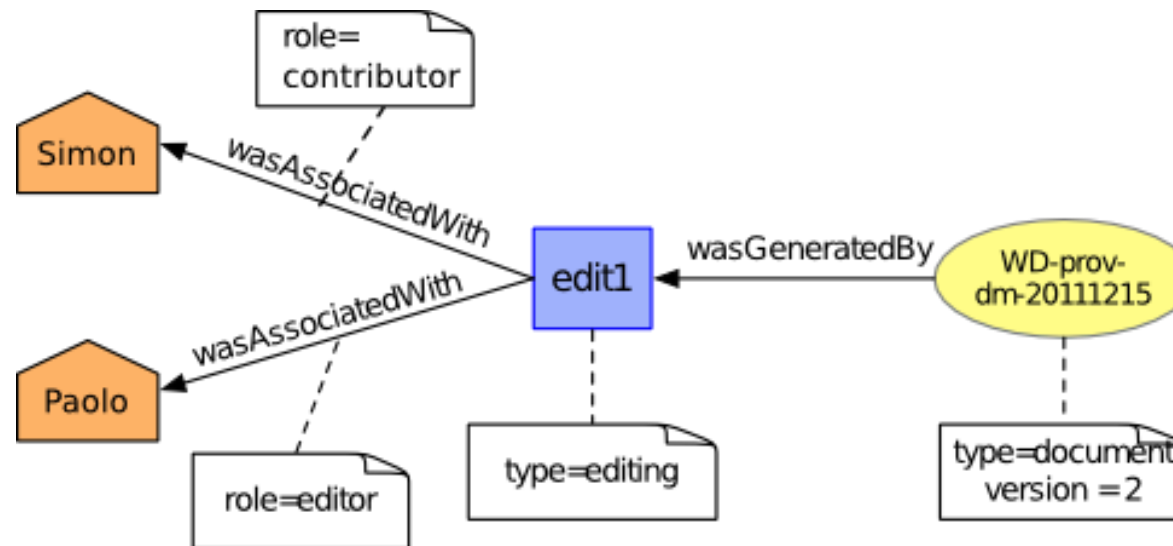
Are we done? No!

- Lacking smart support to govern the complexity of data and transformations
- Data transformations must be governed to prevent DP turns into a swamp
 - Amplified in data science, with data scientists prevailing data architects
 - Leverage descriptive metadata and maintenance to keep control over data

Data platform

Data provenance

- Metadata pertaining to the history of a data item
- Pipeline including the origin of objects and operations they are subjected to
- We have a standard: <https://www.w3.org/TR/prov-dm/>



<https://www.w3.org/TR/prov-dm/>

Data provenance

Entity

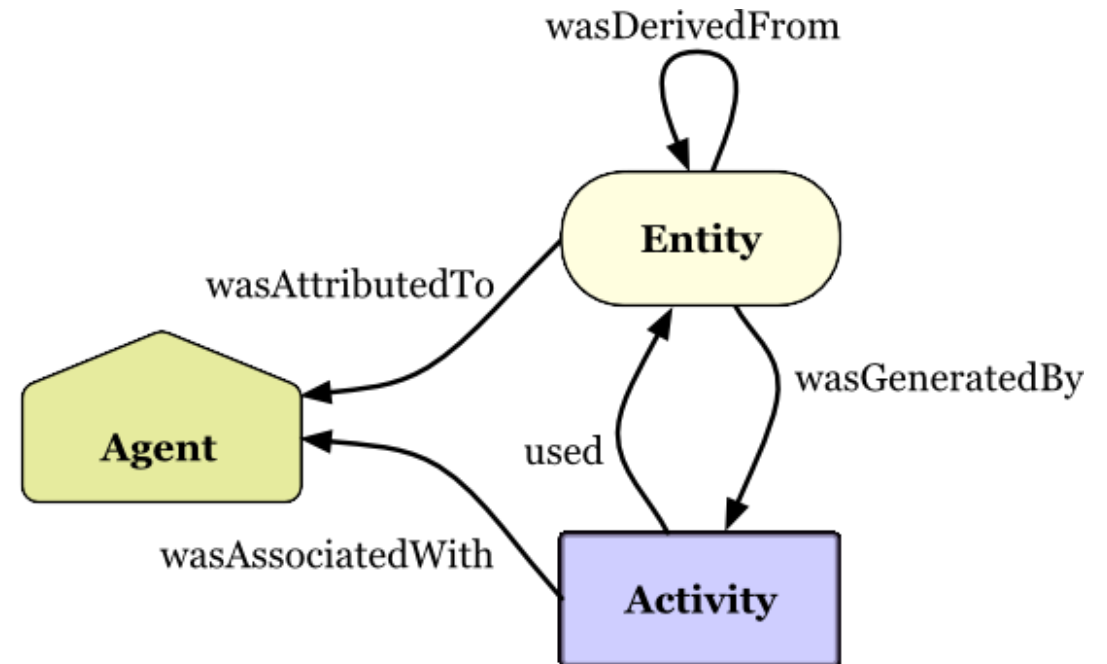
- Physical/conceptual things

Activity

- Dynamic aspects of the world, such as actions
- How entities come into existence, often making use of previously existing entities

Agent

- A person, a piece of software
- Takes a role in an activity such that the agent can be assigned some degree of responsibility for the activity taking place



<https://www.w3.org/TR/2013/NOTE-prov-primer-20130430/>

Data provenance

Data quality

- Monitoring of the quality (e.g., accuracy) of the objects produced
- Notify when a transformation pipeline is not behaving as expected

Debugging

- Inferring the cause of pipeline failures is challenging
- Store inputs of each operation with versions and environmental settings (RAM, CPUs, etc.)

And so on...

Data platform

Are we done? No!

- Metadata can become bigger than data themselves

We need meta meta-data (or models)...

- ... chasing our own tails

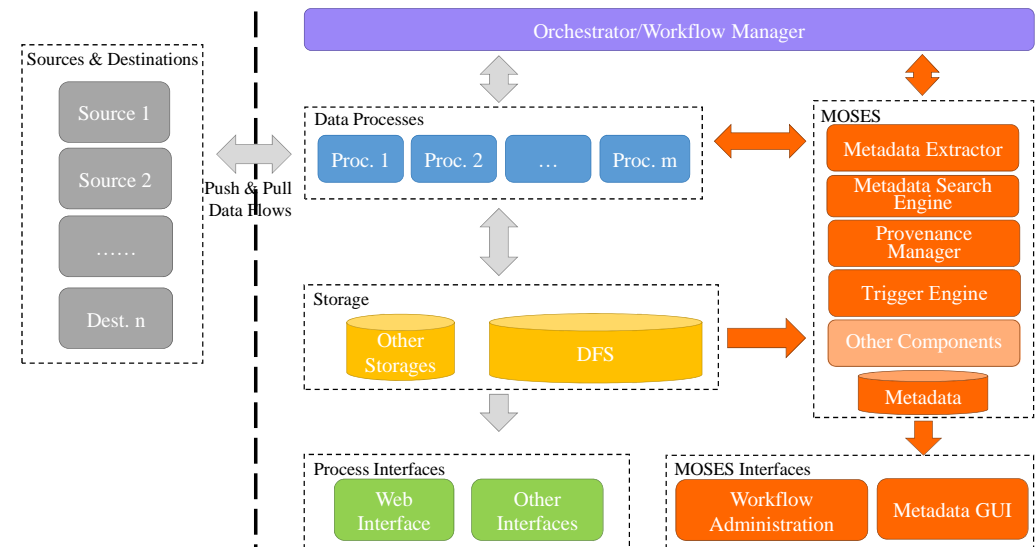
Data management is still a (research) issue in data platforms

Data platform

MOSES: making data platform smarter

Functional architecture

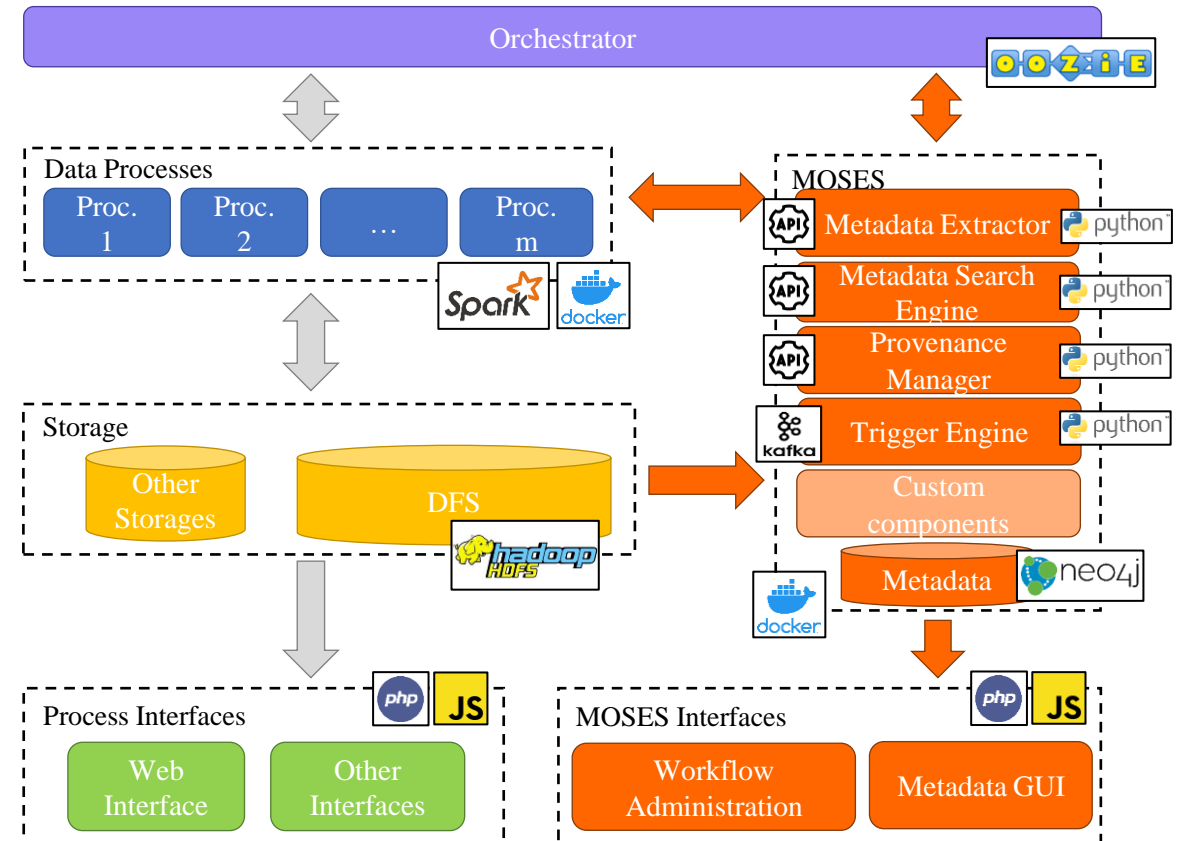
- Components of MOSES are in orange
- The others are standard components in charge of producing/consuming, processing, storing, and visualizing data
- The orchestrator (e.g., Oozie in the Hadoop ecosystem) manages (e.g., schedules) the data transformation processes



Data platform

MOSES is used within a DP that builds atop the Hadoop ecosystem


- The cluster runs the Cloudera Distribution for Apache Hadoop 6.2.0 and Docker
- The metadata collected and manipulated are stored in a graph database
 - The graph data model favors the modeling of highly interconnected data in the absence of fixed schemas
- Service decoupling: functional components are accessible using RESTful APIs



Leoni Anna Giulia. "Gestione di un data lake strutturato attraverso il riconoscimento semantico dei dati acquisiti."

BIG DATA

From PaaS to FaaS (serverless)




<https://xkcd.com/1084/>

From PaaS to FaaS

Public: accessible to anyone willing to pay (e.g., Microsoft, AWS, Google)

Private: accessible by individuals within an institution

Hybrid: a mix of the previous



From PaaS to FaaS

Figure 1. Magic Quadrant for Cloud Infrastructure and Platform Services



<https://www.gartner.com/en/research/methodologies/magic-quadrants-research>

Gartner Magic Quadrant

- Understanding the technology providers to consider for an investment
- **Leaders** execute well and are well positioned for tomorrow
- **Visionaries** understand where the market is going but do not yet execute well
- **Niche Players** focus successfully on a small segment, or are unfocused and do not out-innovate or outperform others
- **Challengers** execute well but do not demonstrate an understanding of market direction
- Focusing on leaders isn't always the best
 - A niche player may support needs better than a market leader. It depends on how the provider aligns with business goals

From PaaS to FaaS

Cloud services are hosted in multiple locations worldwide

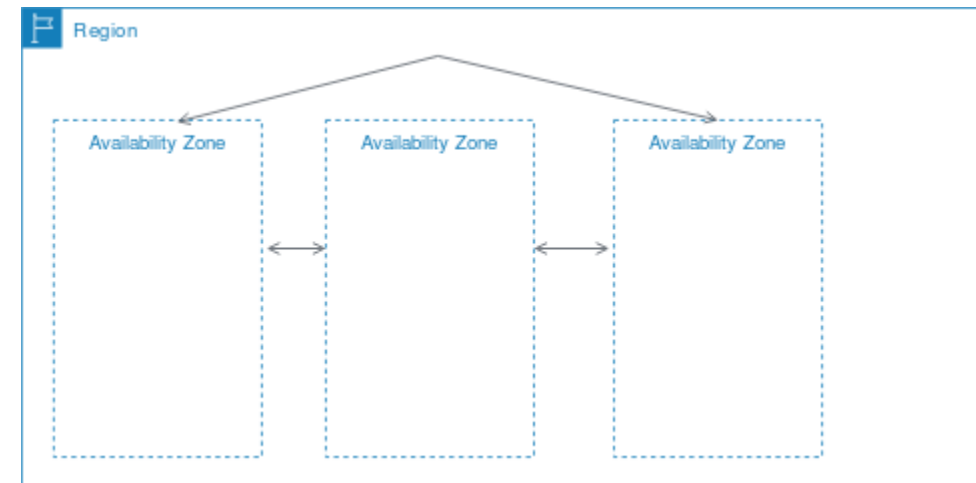
- Locations are composed of **regions** and **availability zones**

Region (e.g., us-east-1)

- Is an independent geographical area
- Has availability zones

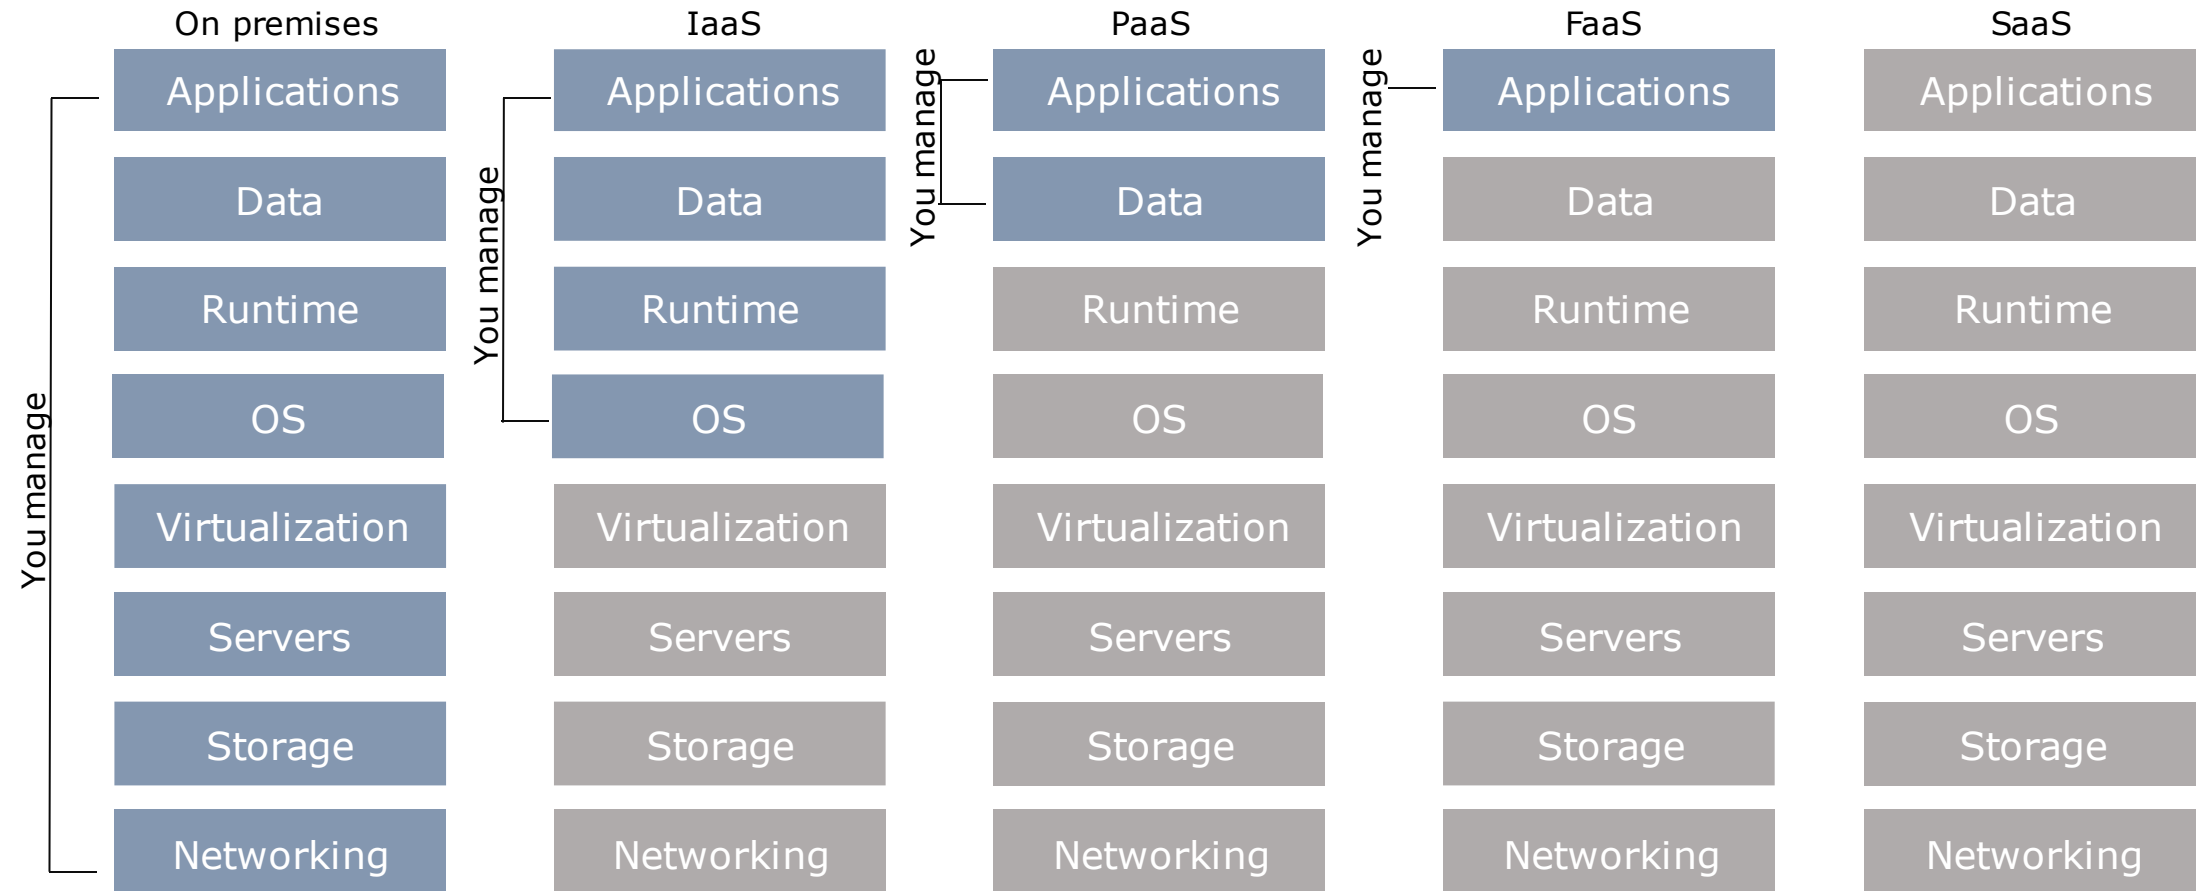
Availability zones in a region

- Are connected through low-latency links
- Resources are usually replicated across zones but not regions



<https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/using-regions-availability-zones.html>

From PaaS to FaaS



From PaaS to FaaS

Understanding architectures is paramount to successful software systems

- Good architectures help to scale
- Poor architectures cause issues that necessitate a costly rewrite

On-premises

- Provisioning, managing, and patching servers is time-consuming
- Require dedicated operations people
- A non-trivial environment is hard to set up and operate effectively
- Infrastructure and hardware are often a distraction from strategic tasks

From PaaS to FaaS

Infrastructure as a service (IaaS)

- A computing infrastructure provisioned and managed over the internet (e.g., AWS EC2)
- Avoid expense/complexity of buying/managing physical servers/data-centers
- IaaS overcomes issues on-premises
- Possibly requires to manage many environments

Platform as a Service (PaaS)

- A development and deployment environment in the cloud (e.g., AWS Elastic Beanstalk)
- Support complete application life-cycle: building, testing, deploying, etc.
- Avoid expense/complexity of managing licenses and application infrastructure

From PaaS to FaaS

PaaS and **containers** are potential solutions to inconsistent infrastructures

PaaS provides a platform for users to run their software

- Developers write software targeting features/capabilities of the platform

Containerization isolates an application with its own environment

- Lightweight alternative to full virtualization
- Containers are isolated but need to be deployed to (public/private) server
- Excellent solution when dependencies are in play
- Housekeeping challenges and complexities

From PaaS to FaaS

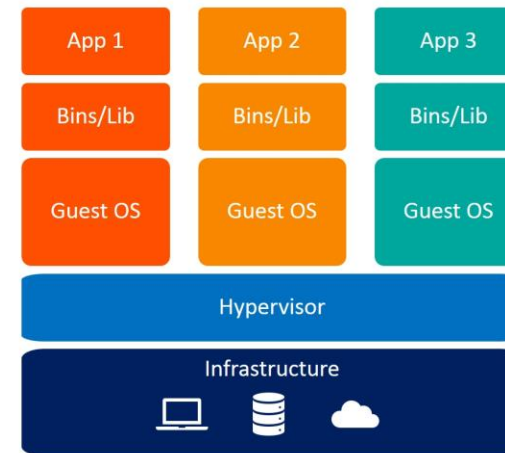
Containers and virtual machines are packaged computing environments

Containers

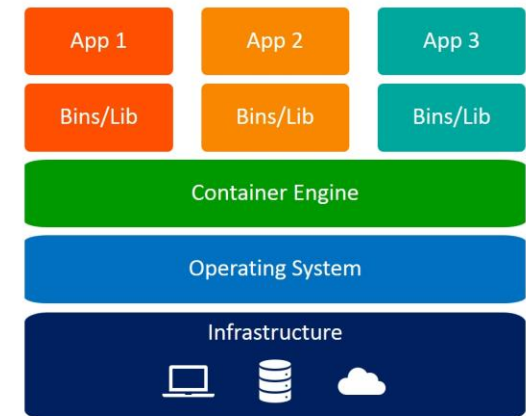
- On top of physical server and its host OS
- Share the host OS kernel
- Shared components are read-only
- Containers are “light” and take just seconds to start

Virtual machines

- Emulate a hardware/software system
- On top of a hypervisor (VM monitor)



Virtual Machines



Containers

From PaaS to FaaS

Function as a Service (FaaS)

- A coding environment, cloud provider provisions platform to run the code (e.g., AWS Lambda)
- Infrastructure provisioning and management are invisible to the developer

Software as a service (SaaS)

- An application environment
- Access cloud-based apps over the Internet (e.g., email, Microsoft Office 365)

From PaaS to FaaS

Serverless

- A software architecture that does not rely on direct access to a server
- Embodies principles from microservices
 - Small, standalone, fully independent services built for a specific purpose
 - Cloud provider is responsible for integration

Principles of FaaS architectures

- FaaS is based on a serverless approach
 - Use a compute service to execute code on demand (no servers/containers)
- Every function could be considered as a standalone service
- Write single-purpose stateless functions

From PaaS to FaaS

Functions react to events

- Design push-based, event-driven pipelines
- Create thicker, more powerful front ends
- Embrace third-party services (e.g., security)

FaaS is not a silver bullet

- Not appropriate for latency-sensitive applications
- Strict specific service-level agreements
- Migration costs
- Vendor lock-in can be an issue