

# NoSQL Databases

---

# About me

## Matteo Francia, PhD

- Adjunct Professor, will teach *Big Data and Cloud Platforms* next year
- Member of the [Business Intelligence Group](#)
- Main research interests:
  - big data
  - data modeling
  - precision agriculture

## Acknowledgements

- Thanks to Enrico Gallinucci for this teaching material

# About today

## A quick introduction on NoSQL databases

- Origins
- Characteristics
- Data models
- Issues
- Usage
- Demo

# Let's start from the beginning

# RDBMSs are full of strength

## ACID properties

- Provides guarantees in terms of consistency and concurrent accesses

## Data integration and normalization of schemas

- Several application can share and reuse the same information

## Standard model and query language

- The relational model and SQL are very well-known standards
- The same theoretical background is shared by the different implementations

## Robustness

- Have been used for over 40 years

# RDBMSs have weaknesses as well

## Impedance mismatch

- Data are stored according to the relational model, but applications to modify them typically rely on the object-oriented model
- Many solutions, no standard
  - E.g.: Object Oriented DBMS (OODBMS), Object-Relational DBMS (ORDBMS), Object-Relational Mapping (ORM) frameworks

## Painful scaling-out

- Not suited for a cluster architecture
- Distributing an RDBMS is neither easy nor cheap (e.g., Oracle RAC)

## Consistency vs latency

- Consistency is a must – even at the expense of latency
- Today's applications require high reading/writing throughput with low latency

## Schema rigidity

- Schema evolution is often expensive

# What is "NoSQL"

The term has been first used in '98 by Carlo Strozzi

- It referred to an open-source RDBMS that used a query language different from SQL

In 2009 it was adopted by a meetup in San Francisco

- Goal: discuss open-source projects related to the newest databases from Google and Amazon
- Participants: Voldemort, Cassandra, Dynamite, HBase, Hypertable, CouchDB, MongoDB

Today, **NoSQL** indicates **DBMSs** adopting a **different data model from the relational one**

- NoSQL = Not Only SQL
- According to Strozzi himself, NoREL would have been a more proper noun

# The first NoSQL systems

## LiveJournal, 2003

- Goal: reduce the number of queries on a DB from a pool of web servers
- Solution: [Memcached](#), designed to keep queries and results in RAM

## Google, 2005

- Goal: handle Big Data (web indexing, Maps, Gmail, etc.)
- Solution: [BigTable](#), designed for scalability and high performance on Petabytes of data

## Amazon, 2007

- Goal: ensure availability and reliability of its e-commerce service 24/7
- Solution: [DynamoDB](#), characterized by strong simplicity for data storage and manipulation



# NoSQL common features

## Not just rows and tables

- Several data model adopted to store and manipulate data

## Freedom from joins

- Joins are either not supported or discouraged

## Freedom from rigid schemas

- Data can be stored or queried without pre-defining a schema (*schemaless* or *soft-schema*)

## Distributed, shared-nothing architecture

- Trivial scalability in a distributed environment with no performance decay
- Each workstation uses its own disks and RAM

# NoSQL misconceptions

## Not a farewell to SQL

- Some systems do adopt SQL (or a SQL-like language)

## Not necessarily open-source

- There exist both open-source and commercial systems

## Not only Cloud Computing

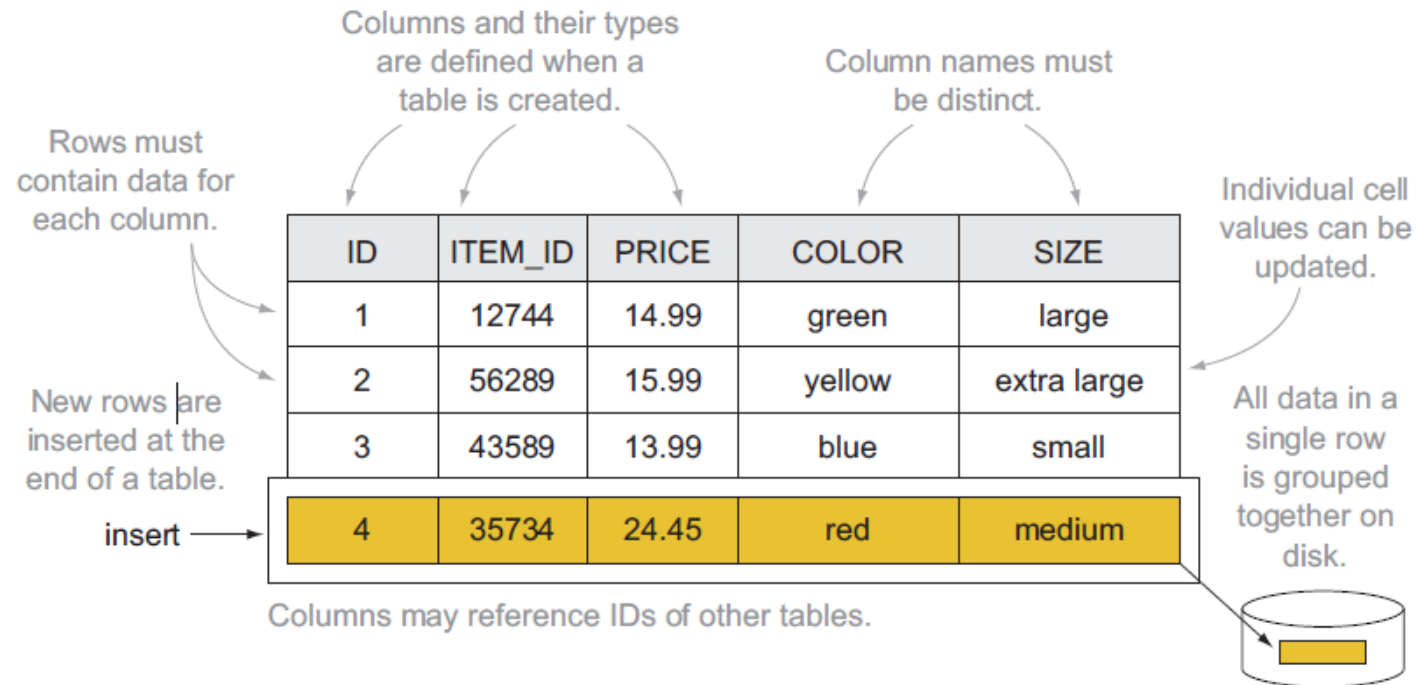
- There exist both on-premise and cloud solutions

## Not an optimization of hardware resources

- With the same resources, a centralized RDBMS is probably better performing

# Data modeling

# Relational model



# NoSQL: several data models

One of the key challenges is to understand which one fits best with the required application

Model	Description	Use cases
Key-value	Associates any kind of value to a string	Dictionary, lookup table, cache, file and images storage
Document	Stores hierarchical data in a tree-like structure	Documents, anything that fits into a hierarchical structure
Wide column	Stores sparse matrixes where a cell is identified by the row and column keys	Crawling, high-variability systems, sparse matrixes
Graph	Stores vertices and arches	Social network queries, inference, pattern matching

# Key-value: data model

Each DB contains one or more **collections** (corresponding to tables)

Each collection contains a list of **key-value pairs**

- Key: a unique string
  - E.g.: ids, hashes, paths, queries, REST calls
- Value: a BLOB (binary large object)
  - E.g.: text, documents, web pages, multimedia files

**Atomicity level:** the key-value pair

Key	Value
image-12345.jpg	Binary image file
http://www.example.com/my-web-page.html	HTML of a web page
N:/folder/subfolder/myfile.pdf	PDF document
9e107d9d372bb6826bd81d3542a419d6	The quick brown fox jumps over the lazy dog
view-person?person-id=12345&format=xml	<Person><id> 12345</id>.</Person>
SELECT PERSON FROM PEOPLE WHERE PID="12345"	<Person><id> 12345</id>.</Person>

Looks like a simple dictionary

- The collection is indexed by key
- The value may contain several information: one or more definitions, synonyms and antonyms, images, etc.

# Key-value: querying

Three simple kinds of query:

- `put($key as xs:string, $value as item())`
  - Adds a key-value pair to the collection
  - If the key already exists, the value is replaced
- `get($key as xs:string) as item()`
  - Returns the value corresponding to the key (if it exists)
- `delete($key as xs:string)`
  - Deletes the key-value pair

The value is a *black box*: it cannot be queried!

- No "where" clauses
- No indexes on the values
- Schema information is often indicated in the key

Key	Value
user:1234:name	Matteo
user:1234:age	33
post:9876:written-by	user:1234
post:9876:title	NoSQL Databases
comment:5050:reply-to	post:9876

# Document: data model

Each DB contains one or more **collections** (corresponding to tables)

Each collection contains a list of **documents** (usually JSON)

- Documents are hierarchically structured

Each document contains a set of **fields**

- The **ID** is mandatory

Each field corresponds to a **key-value pair**

- Key: unique string in the document
- Value: either simple (string, number, boolean) or complex (object, array, BLOB)
  - A complex field can contain other field

**Atomicity level**: the document

```
{
  "_id": 1234,
  "name": "Matteo",
  "age": 33,
  "address": {
    "city": "Ravenna",
    "postalCode": 48124
  },
  "contacts": [ {
    "type": "office",
    "contact": "0547-338835"
  }, {
    "type": "skype",
    "contact": "mfrancia"
  } ]
}
```



# Document: querying

Differently from the key-value, the value is *visible* by the DBMS

Thus, query languages are quite expressive

- Can create indexes on fields
- Can filter on the fields
- Can return more documents with one query
- Can select which fields to project
- Can update specific fields

Different implementations, different functionalities

- Some enable (possibly materialized) views
- Some enable MapReduce queries
- Some provide connectors to Big Data tools (e.g., Spark, Hive)
- Some provide *full-text search* capabilities

# Wide column: data model

Each DB contains one or more **column families** (corresponding to tables)

Each column family contains a list of **row** in the form of a key-value pair

- Key: unique string in the column family
- Value: a set of **columns**

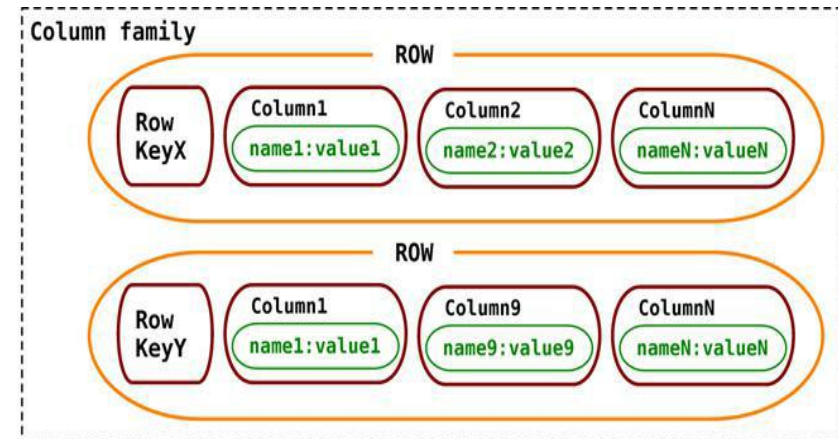
Each column is a key-value pair itself

- Key: unique string in the row
- Value: simple or complex (*supercolumn*)

**Atomicity level:** the row

With respect to the relational model:

- Rows specify only the columns for which a value exists
  - Particularly suited for sparse matrixes
- Timestamps can be used to defines *versions* of column values



Column family	Row key	Column key	Timestamp	Value
---------------	---------	------------	-----------	-------

# Wide column: querying

The query language expressiveness is in between key-value and document data models

- Column indexes are discouraged
- Can filter on column values (not always)
- Can return more rows with one query
- Can select which columns to project
- Can update specific columns (not always)

Given the similarity with the relational model, a **SQL-like** language is often used

# Wide column: $\neq$ columnar

Do not mistake the wide column data model with the columnar storage used for OLAP applications

A	B	C
A1	B1	C1
A2	B2	C2
A3	B3	C3



Row-oriented

Column-oriented

## Row-oriented

- Pro: inserting a record is easy
- Con: several unnecessary data may be accessed when reading a record

## Column-oriented

- Pro: only the required values are accessed
- Con: writing a record requires multiple accesses

# Graph: data model

Each DB contains one or more **graphs**


Each graph contains **vertices** and **arcs**

- Vertices: usually represent real-world entities
  - E.g.: people, organizations, web pages, workstations, cells, books, etc.
- Arcs: represent directed relationships between the vertices
  - E.g.: friendship, work relationship, hyperlink, ethernet links, copyright, etc.
- Vertices and arcs are described by **properties**
- Arcs are stored as physical pointers

**Atomicity level:** the transaction

Most known specializations:

- Reticular data model (Parent-child or owner-member relationships)
- Triplestore (Subject-predicate-object relationships; e.g., RDF)



# Graph: querying

Graph databases usually model completely different contexts

Thus, query language and mechanism is quite different

- Support for transactions
- Support for indexes, selections and projections
- Query language based on detecting patterns

Query	Pattern
Find friends of friends	(user)-[:KNOWS]-(friend)-[:KNOWS]-(foaf)
Find shortest path from A to B	shortestPath( (userA)-[:KNOWS*..5]-(userB) )
What has been bought by those who bought my same products?	(user)-[:PURCHASED]->(product)<-[:PURCHASED]-()-[:PURCHASED]->(otherProduct)

# Aggregate vs Graph modeling

Key-value, document and wide column are called **aggregate-oriented**

- Aggregate = key-value pair, document, row (respectively)
- The aggregate is the atomic block (no guarantees for multi-aggregate operations)

Based on the concept of encapsulation

- Pro: avoid joins as much as possible → achieve **high scalability**
- Con: data denormalization → **potential inconsistencies in the data**
- **Query-driven modeling**

The graph data model is intrinsically different from the others

- Focused on the relationships rather than on the entities per-se
- **Limited scalability**: it is often impossible to shard a graph on several machines without "cutting" several arcs (i.e. having several cross-machine links)
  - Batch cross-machine queries: don't follow relationships one by one, but "group them" to make less requests
  - Limit the depth of cross-machine node searches
- **Data-driven modeling**

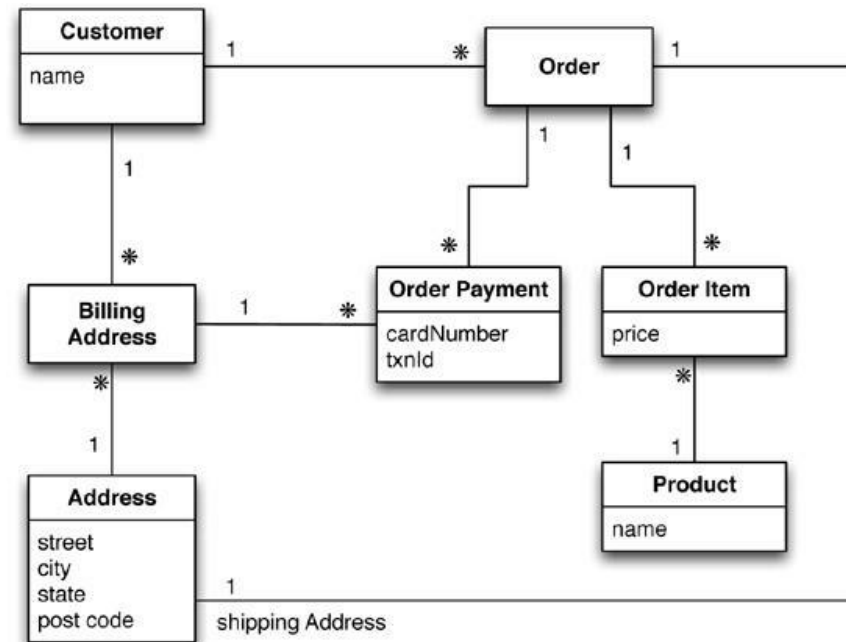
# Data modeling

Let's see some examples



# Data modeling

Typical use case: customers, orders and products



# Data modeling: relational

Customer	
Id	Name
1	Martin

Orders		
Id	CustomerId	ShippingAddressId
99	1	77

Product	
Id	Name
27	NoSQL Distilled

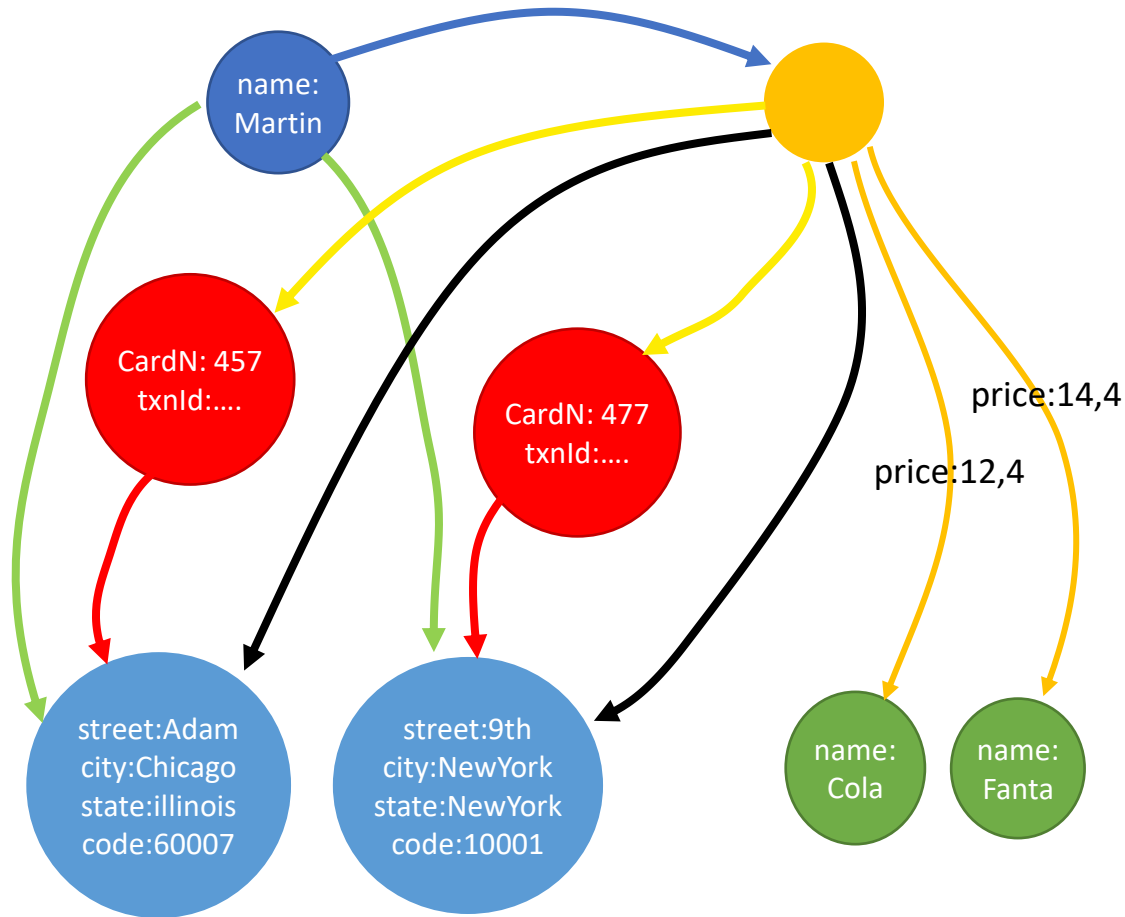
BillingAddress		
Id	CustomerId	AddressId
55	1	77

OrderItem			
Id	OrderId	ProductId	Price
100	99	27	32.45


Address	
Id	City
77	Chicago

OrderPayment				
Id	OrderId	CardNumber	BillingAddressId	txnId
33	99	1000-1000	55	abelif879rft

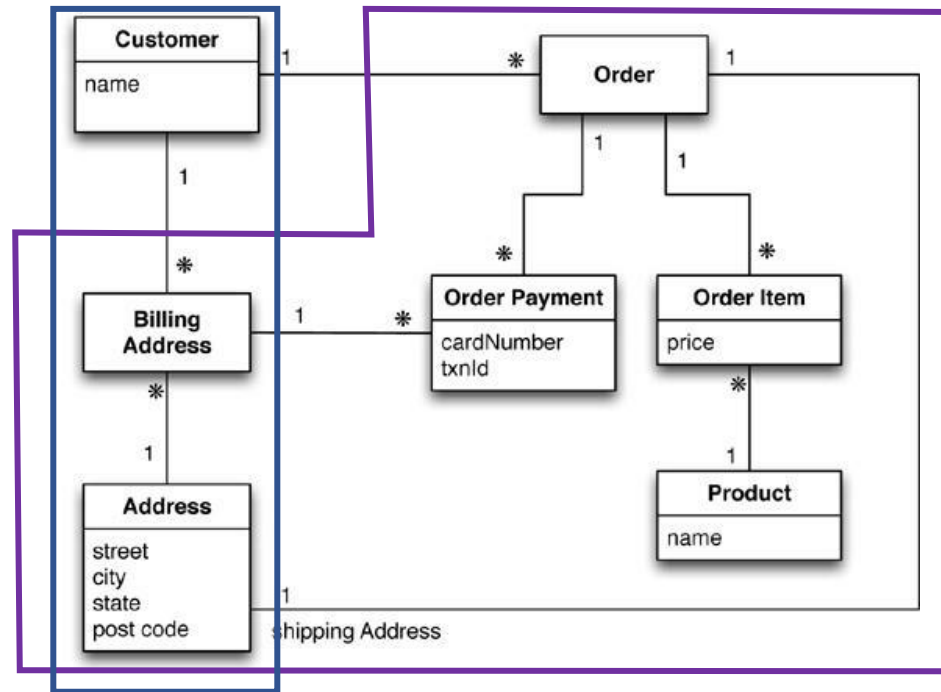
# Data modeling: graph



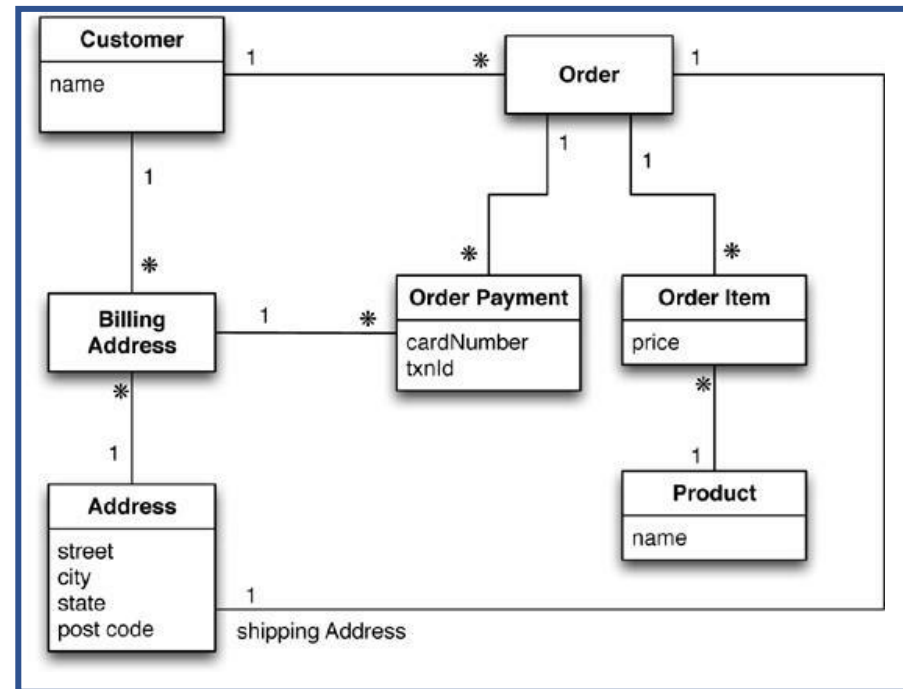
- IDs are implicit
- Different edge colors imply different edge types



# Data modeling: aggregate-oriented



# Data modeling: aggregate-oriented



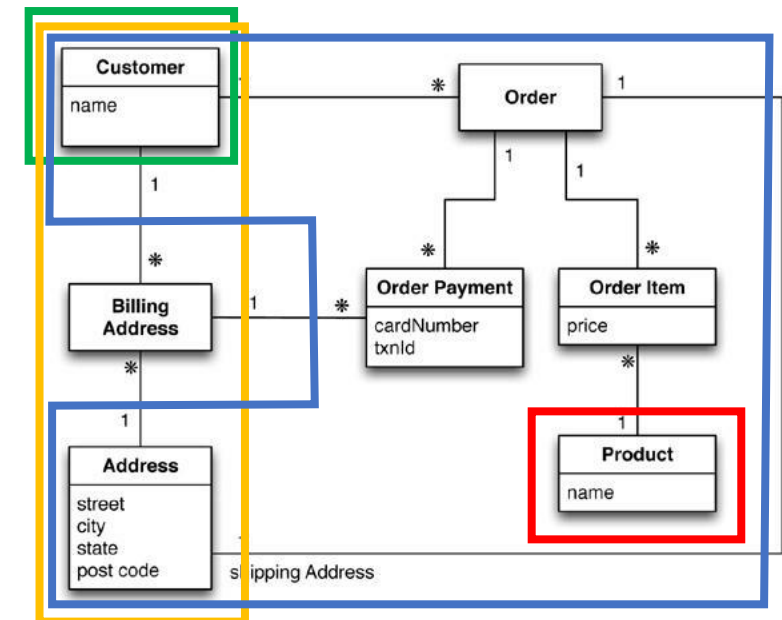
# Data modeling: key-value

*Customer collection*

key	value
cust-1:name	Martin
cust-1:adrs	[ {"street":"Adam", "city":"Chicago", "state":"Illinois", "code":60007}, {"street":"9th", "city":"NewYork", "state":"NewYork", "code":10001} ]
cust-1:ord-99	{ "orderpayments": [ {"card":477, "billadrs": {"street":"Adam", "city":"Chicago", "state":"illinois", "code":60007} }, {"card":457, "billadrs": {"street":"9th", "city":"NewYork", "state":"NewYork", "code":10001} }, "products": [ {"id":1, "name":"Cola", "price":12.4}, {"id":2, "name":"Fanta", "price":14.4} ], "shipAdrs": {"street":"9th", "city":"NewYork", "state":"NewYork", "code":10001} }

*Product collection*

key	value
p-1:name	Cola
p-2:name	Fanta



# Data modeling: document-based

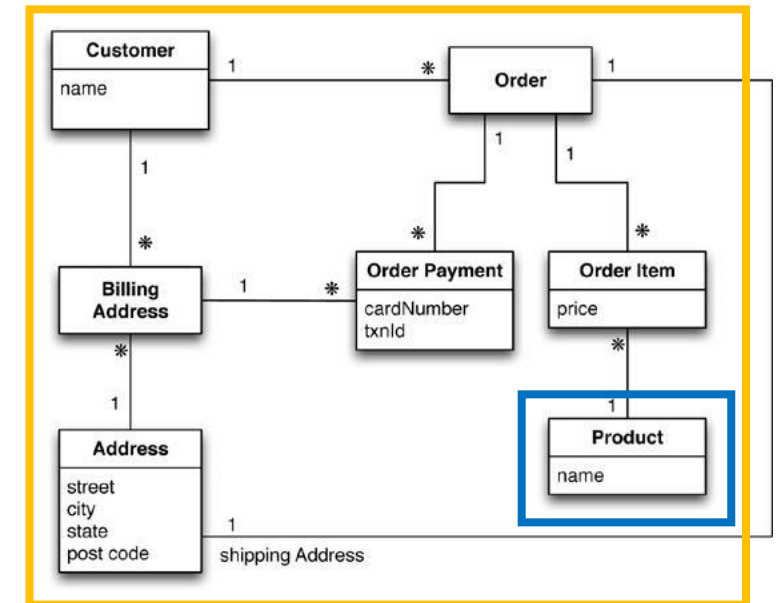
*Customer collection*

```
{
  "_id": 1,
  "name": "Martin",
  "adrs": [
    {"street": "Adam", "city": "Chicago", "state": "illinois", "code": 60007},
    {"street": "9th", "city": "NewYork", "state": "NewYork", "code": 10001}
  ],
  "orders": [ {
    "orderpayments": [
      {"card": 477, "billadrs": {"street": "Adam", "city": "Chicago", "state": "illinois", "code": 60007}},
      {"card": 457, "billadrs": {"street": "9th", "city": "NewYork", "state": "NewYork", "code": 10001}}
    ],
    "products": [
      {"id": 1, "name": "Cola", "price": 12.4},
      {"id": 2, "name": "Fanta", "price": 14.4}
    ],
    "shipAdrs": {"street": "9th", "city": "NewYork", "state": "NewYork", "code": 10001}
  }
}
```

*Product collection*

```
{
  "_id": 1,
  "name": "Cola",
  "price": 12.4
}

{
  "_id": 1,
  "name": "Fanta",
  "price": 14.4
}
```



# Data modeling: document-based

```
{
  "_id": 1,
  "name": "Martin",
  "adrs": [
    {"street": "Adam", "city": "Chicago", "state": "illinois", "code": 60007},
    {"street": "9th", "city": "NewYork", "state": "NewYork", "code": 10001}
  ]
}
```

*Customer*  
collection

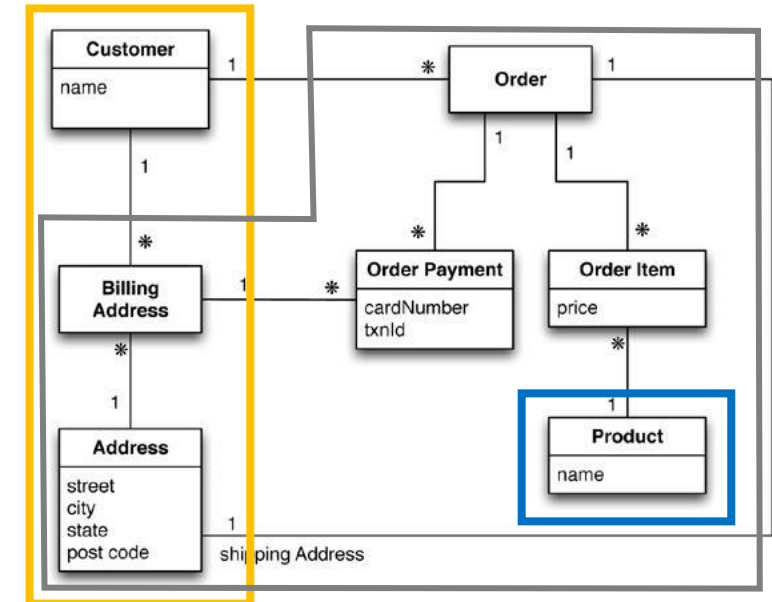
```
{
  "_id": 1,
  "name": "Cola",
  "price": 12.4
}
```

```
{
  "_id": 1,
  "name": "Fanta",
  "price": 14.4
}
```

*Product*  
collection

```
{
  "_id": 1,
  "customer": 1,
  "orderpayments": [
    {"card": 477, "billadrs": {"street": "Adam", "city": "Chicago", "state": "illinois", "code": 60007}},
    {"card": 457, "billadrs": {"street": "9th", "city": "NewYork", "state": "NewYork", "code": 10001}}
  ],
  "products": [
    {"id": 1, "name": "Cola", "price": 12.4},
    {"id": 2, "name": "Fanta", "price": 14.4}
  ],
  "shipAdrs": {"street": "9th", "city": "NewYork", "state": "NewYork", "code": 10001}
}
```

*Order*  
collection





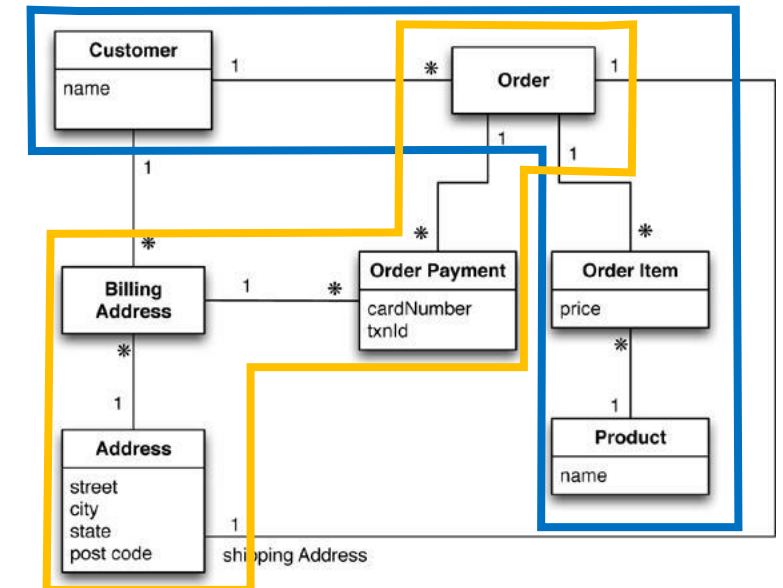
# Data modeling: wide-column

*Order details* column family

Ord	CustName	Pepsi	Cola	Fanta	...
1	Martin		12.4	14.4	
2	...	...			...

*Order payments* column family

Ord	OrderPayments				
1	Card	Steet	City	State	Code
	477	9th	NewYork	NewYork	10001
	457	Adam	Chicago	Illinois	60007
2	...				



# Summary

The *aggregate* term comes from Domain-Driven Design

- An aggregate is a group of tightly coupled objects to be handled as a block
- Aggregates are the basic unit for data manipulation and consistency management

## Advantages

- **Can be distributed trivially** (group the data that are used altogether)
- Facilitate the developer's job (by surpassing the impedance mismatch problem)

## Disadvantages

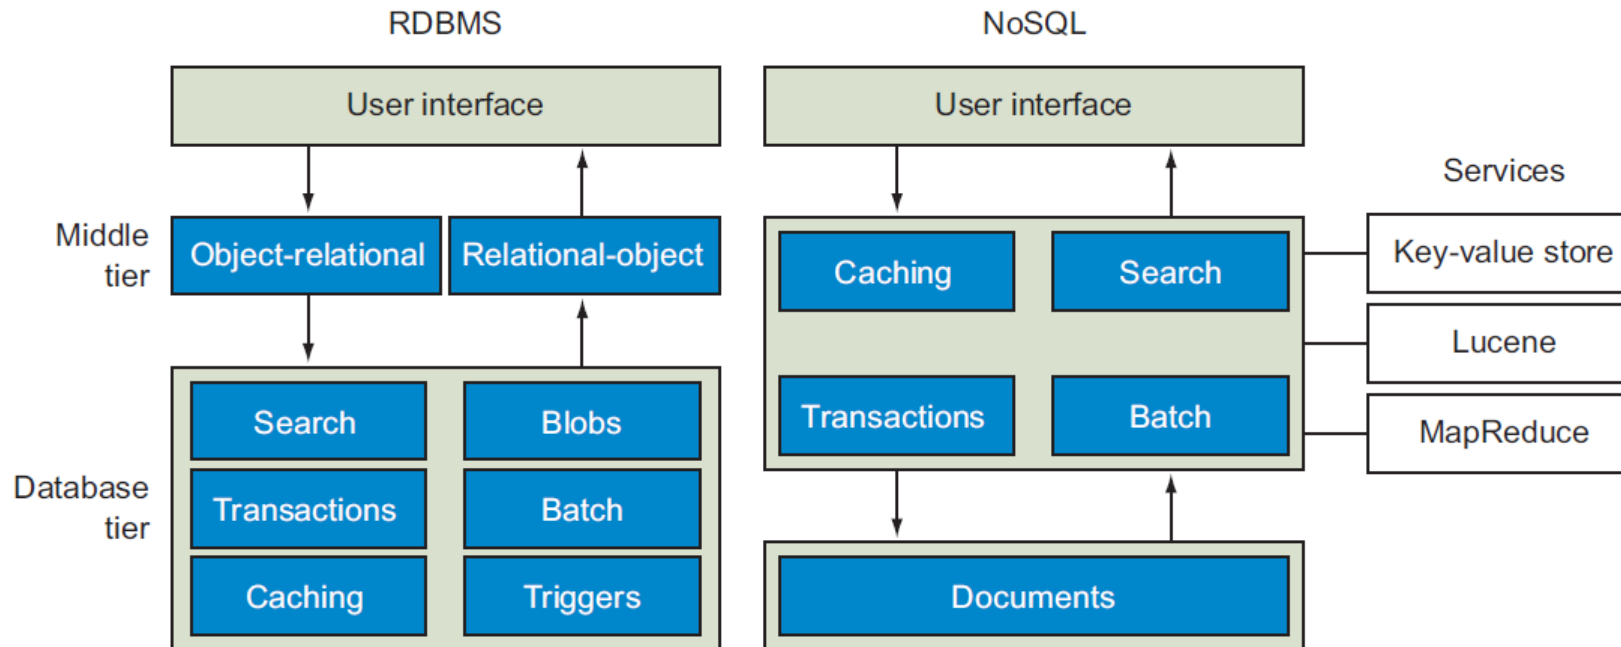
- **No real design strategy exists for aggregates**
  - It mainly depends on how data are meant to be used
- Can optimize only a limited set of queries
- Data denormalization → possible inconsistencies

A data-driven approach is agnostic to how the data is used (mostly)

# Dealing with data consistency

In short

# RDBMS vs NoSQL: different philosophies



# Consistency: an example

Consider 1000€ to be transferred from bank account A to B; the transfer is made by:

- Removing 1000€ from A
- Adding 1000€ to B

What should never happen

- The money is removed from A but not added to B
- The money is added twice to B
- A query on the database shows an intermediate state
  - E.g.,  $A+B = 0\text{€}$

RDBMS adopt **transactions** to avoid this kind of issue

# Consistency in NoSQL: CAP

"Theorem": only two of the following three properties can be guaranteed

**Consistency:** the system is always consistent


- Every node returns the same, most recent, successful write
- Every client has the same view of the data

**Availability:** the system is always available

- Every non-failing node returns a response for all read and write requests in a reasonable amount of time

**Partition tolerance:** the system continues to function and upholds its consistency guarantees in spite of network partitions

- In distributed systems, network partitioning is an *inevitable possibility*



# Consistency in NoSQL: CAP

## Three situations

- CA: the system cannot suffer from network partitioning (single server)
- AP: in case of partitioning, the system sacrifices consistency (overbooking)
- CP: in case of partitioning, the system sacrifices availability (bookings prevented)

## Theorem interpretation is not trivial

- Asymmetric properties: consistency is sacrificed to favor speed at all times, not just when partitioning happens
- Different application requirements → different algorithms handle these properties more strictly/loosely

# Consistency in NoSQL: CAP relaxed

Consider two users that want to book the same room when a network partition happens

**CP:** no one can book (**A is sacrificed**)

- Not the best solution

**AP:** both can book (**C is sacrificed**)

- Possible overbooking: writing conflict to handle

**caP:** only one can book

- The other will see the room available but cannot book it

**This is admissible only in certain scenarios**

- Finance? Blogs? E-commerce?

It's important to understand:

- What is the tolerance to obsolete reads
- How large can the inconsistency window be



# Consistency in NoSQL: summary

NoSQL databases rely on the *eventual consistency* principle

- Trade absolute consistency for latency
- Inconsistencies are tolerated within short time windows

Still, *unrecoverable inconsistencies* may arise

- Because of data denormalization
- Because of the absence of transaction

One size does not fit all

# Key-Value: popular DBs

**Redis** (Data Structure server): <http://redis.io/>

- Supports complex fields (list, set, ...) and operations on values (range, diff, ...)

**Memcached DB:** <http://memcached.org/>

**Riak:** <http://basho.com/riak/>

# Key-Value: when to use

## Very simple use cases

- Independent data (no need to model relationships)
- The typical query is a simple lookup
- Need super-fast performance

## Examples

- **Session information**
  - Each web session is identified by its own sessionId: All related data can be stored with a `PUT` request and returned with a `GET` request.
- **User profiles, preferences**
  - Each user is uniquely identified (userId, username) and has her own preferences in terms of language, colors, timezone, products, etc. – data that fits well within an aggregate
- **Shopping cart, chat services**
  - Each e-commerce websites associates a shopping cart to a user; it can be stored as an aggregate identified by the user ID.

# Key-Value: real use cases

## Crawling of web pages

- The URL is the key, the whole page content (HTML, CSS, JS, images, ..) is the value

Key	Value
http://www.example.com/index.html	<html>...
http://www.example.com/about.html	<html>...
http://www.example.com/products.html	<html>...
http://www.example.com/logo.png	Binary...

## Twitter timeline

- The user ID is the key, the list of most recent tweets to be shown is the value

## Amazon S3 (Simple Storage Service)

- A cloud-based file system service
- Useful for personal backups, file sharing, website or apps publication
- The more you store, the more you pay
  - Storage: approx. \$0.03 per GB per month
  - Uploading files: approx. \$0.005 per 1000 items
  - Downloading files: approx. \$0.004 per 10,000 files\* PLUS \$0.09 per GB (first GB free)

# Key-Value: when to avoid

## **Data with many relationships**

- When relationships between data (in the same or in different collections) must be followed
- Some systems offer limited link-walking mechanisms

## **Multi-record operations**

- Because operations (mostly) involve one record at a time

## **Querying the data**

- If it is necessary to query the values, not just the key
- Few systems offer limited functionalities (e.g., Riak Search)

# Document: popular DBs

**MongoDB:** <http://www.mongodb.org>

**Couchbase:** <http://www.couchbase.com>

**CouchDB:** <http://couchdb.apache.org>

# Document: when to use

## Higher expressiveness

- Store data according to a highly nested data model
- Need to formulate complex queries on many fields

## Examples

- **Event logs**
  - Central repo to store event logs from many applications; shard on app name or event type
- **CMS, blogging platforms**
  - The absence of a predefined schema fits well within content management systems (CMS) or website management applications, to handle comments, registrations and user profiles
- **Web Analytics or Real-Time Analytics**
  - The ability to update only specific fields enables fast update of analytical metrics
  - Text indexing enables real-time sentiment analysis and social media monitoring
- **E-commerce applications**
  - Schema flexibility is often required to store products and orders, as well as to enable schema evolution without incurring into refactoring or migration costs



# Document: real use cases

## Advertising services

- MongoDB was born as a system for banner ads
  - 24/7 availability and high performance
  - Complex rules to find the right banner based on user's interests
  - Handle several kinds of ads and show detailed analytics

## Internet of Things

- Real-time management of sensor-based data
- Bosch uses MongoDB to capture data from cars (breaks, ABS, windscreen wiper, etc.) and aircrafts maintenance tools
  - Business rules are applied to warn the pilot when the breaking system pressure falls under a critical threshold, or the maintenance operator when the tool is used improperly
- Technogym uses MongoDB to capture data from gym equipment

# Document: when to avoid

## ACID transactions requirement

- If not for a few exceptions (e.g., RavenDB), document databases are not suited for cross-document atomicity

## Queries on high-variety data

- If the aggregate structure continuously evolves, queries must be constantly updated (and normalization clashes with the concept of aggregate)

# Wide column: popular DBs

**Cassandra:** <http://cassandra.apache.org>

**HBase:** <https://hbase.apache.org>

**Google BigTable:** <https://cloud.google.com/bigtable>

# Wide column: when to use

## Compromise between expressiveness and simplicity

- Limited (but some) requirements in terms of data model
- Limited (but some) requirements in terms of querying records

## Examples

- **Event logs; CMS, blogging platforms**
  - Similarly to document databases, [different applications may use different columns](#)
- **Sparse matrixed**
  - While an RDBMS would store *null* values, a wide column [stores only the columns for which a value is specified](#)
- **GIS applications**
  - Pieces of a map (tiles) can be stored as [couples of latitude and longitude](#)

# Wide column: real use cases

## **Google applications**

- BigTable is the DB used by Google for most of its applications, including Search, Analytics, Maps and Gmail

## **User profiles and preferences**

- Spotify uses Cassandra to store metadata about users, artists, songs, playlists, etc.

# Wide column: when to avoid

## **Same as for document model**

- ACID transactions requirement
- Queries on high-variety data

## **Need for full query expressiveness**

- Joins are usually *not* supported
- Limited support for filters and group bys

# Graph: popular DBs

**Neo4J:** <http://neo4j.com>

**TigerGraph:** <https://www.tigergraph.com/>

# Graph: when to use

## Interlinked data

- **Social networks** are one of the most typical use case of graph databases (e.g., to store friendships or work relationships); **every relationship-centric domain is a good one**

## Routing and location-based services

- Applications working on the **TSP (Travelling Salesman Problem)** problem
- Location-based application that, for instance, recommend the best restaurant nearby; in this case, **relationships model the distance between node**

## Recommendation applications, fraud-detection

- Systems recommending «the products bought by your friends», or «the products bought by those who bought your same products»
- When relationships model behaviors, outlier detection may be useful to identify frauds



# Graph: real use cases

## Relationships analysis

- Finding common friends (e.g., friend-of-a-friend) in a social network
- Identifying clusters of phone calls that identify a criminal network
- Analyzing flows of money to identifying money recycling patterns or credit card theft
- Main users: law firms, police, intelligence agencies
  - <https://neo4j.com/use-cases/fraud-detection/>
- Useful for text analysis as well (Natural Language Processing)

## Inference

- Creating rules that define new knowledge based on existing patterns (e.g., transitive relationships, trust mechanisms)

# Graph: when to avoid

## Data-intensive applications

- Traversing the graph is trivial, but [analyzing the whole graph can be expensive](#)
- There exist frameworks for distributed graph analysis (e.g., Apache Giraph), but they do not rely on a graph DB

So... are RDBMSs dead?

# No! Polyglot persistence

Different databases are designed to solve different problems


Using a single DBMS to handle everything (the *one size fits all*)...

- Operational data
- Temporary session information
- Graph traversing
- OLAP analyses
- ...

... usually lead to inefficient solutions

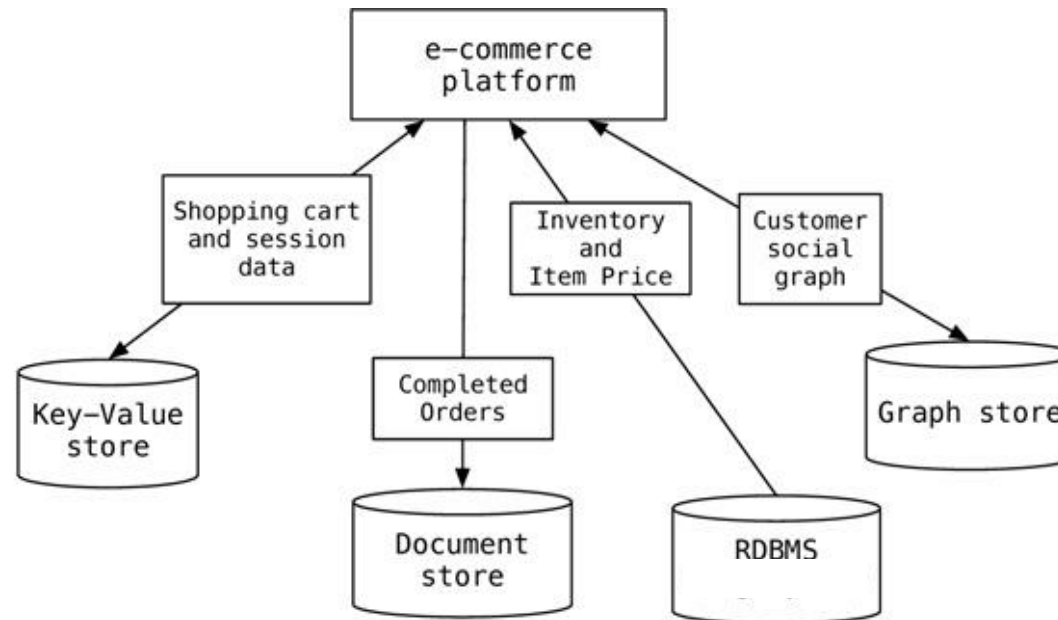
Each activity has its own requirements (availability, consistency, fault tolerance, etc.)

# *One size fits all* approach



# Polyglot data management

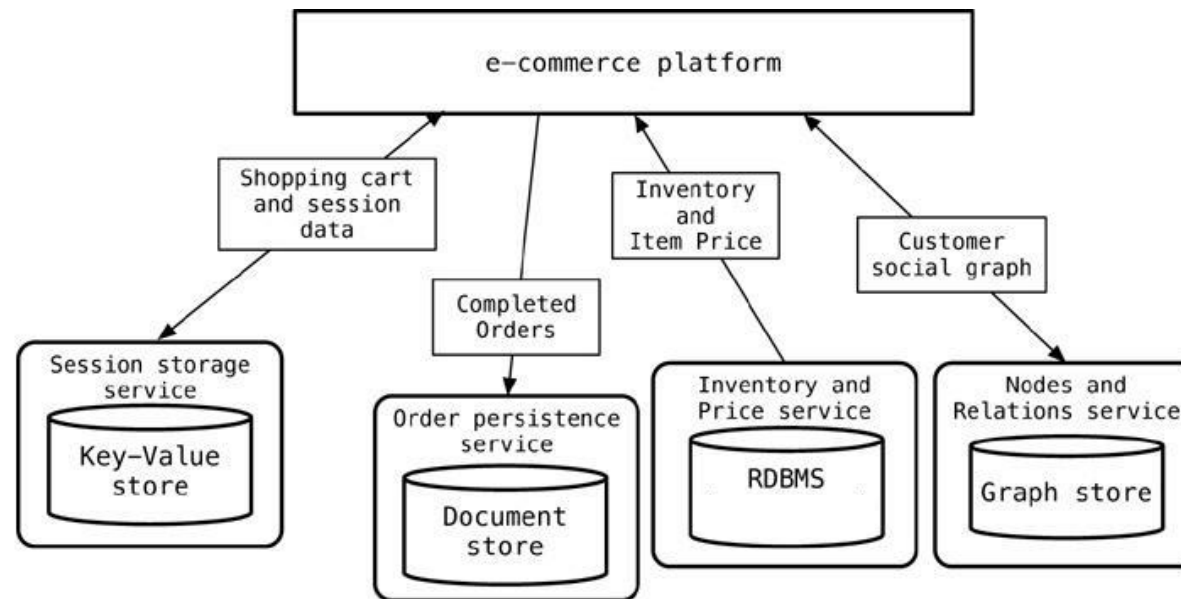
Beware: this leads to harder cross-database consistency management



# Polyglot data management

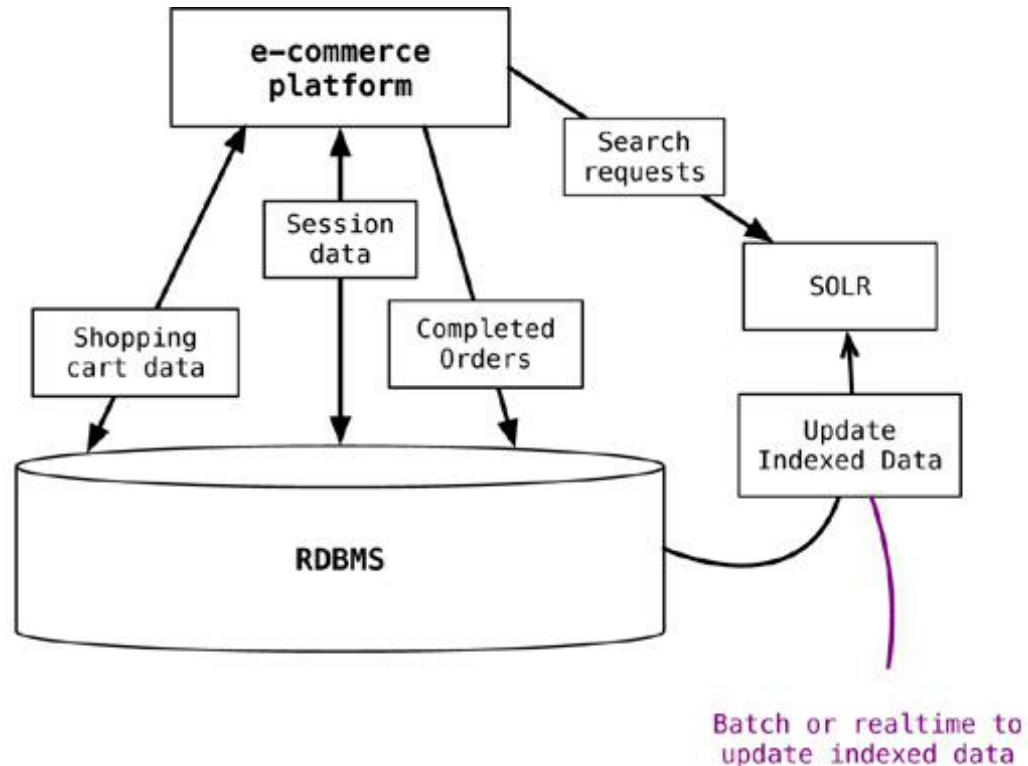
Each DB should be "embedded" within services

- Data access and manipulation enabled with API services



# Supporting existing technologies

If the current solution cannot be changed, NoSQL systems can still support the existing ones





# Beyond NoSQL

## NewSQL systems

- Combine the benefits from both relational and NoSQL worlds
- Ensure scalability without compromising consistency...
- ...but by compromising some availability!

## Multi-model systems

- Support multiple data models
- The barrier between distributed and centralized DBMSs still remains

## Database-as-a-service

- All cloud providers offer storage services supporting all data models

# NoSQL Databases

## Complementary to RDBMSs

- Different data models
- Better scaling
- Trade absolute consistency for latency
- NewSQL and multi-model systems also exist



Want to know more?

← Suggested readings

- [Distributed algorithms in NoSQL databases](#)
  - Email me at [m.francia@unibo.it](mailto:m.francia@unibo.it)
    - See you next year!